



Data Mining in Rust



Chris Pearce
Rust AKL Meetup, 16 Oct 2018



about:me

- Software Engineer, Firefox, since 2007.
- <http://github.com/cpearce>
- Mostly worked on HTML5 <video> in Firefox in C++
- Married to Dr Yun Sing Koh, Senior Lecturer, CS @ University of Auckland
 - Data mining & machine learning expert
 - <https://www.cs.auckland.ac.nz/~yunsing/>
- Association rule mining implementor

Agenda

1. Rust borrow checker rules
2. Describe Association Rule Data Mining
3. Explain FP Growth Algorithm
4. Discuss challenges of using Rust
5. Describe optimizing implementation of FPGrowth
6. Parallelism & concurrency

Rust's Rules of References

The Rust Programming Language, Chapter 4.2

- Each value in Rust has a variable that's called its owner.
- There can only be one owner at a time.
- When the owner goes out of scope, the value will be dropped.
- At any given time, you can have either (but not both of):
 - one mutable reference or,
 - any number of immutable references.
- References must always be valid.

Rules of References; Example

```
let mut s = String::from("hello");
```

```
let r1 = &s; // no problem
```

```
let r2 = &s; // no problem
```

```
let r3 = &mut s; // Error!
```

```
error[E0502]: cannot borrow `s` as mutable because it is also borrowed as  
immutable
```

Association Rule Mining in Rust

<https://github.com/cpearce/arm-rs/>

See also...

<https://github.com/cpearce/arm-py> (Python3)

<https://github.com/cpearce/arm-java> (Java8)

<https://github.com/cpearce/HARM> (ugly C++)

Association Rule Mining:

~~Market basket~~ *Candy Bar* analysis

- Given data set of transactions, find associations between items.

Transaction ID	Items
1	popcorn, coke, choctop, ...
2	wine, jaffas, choctop...
3	coke, crisps, M&Ms, ...
...	...

Two sub problems

1. Generating frequent patterns
2. Generating associations rules

Generating Frequent Patterns

- Input: set of transactions
- Find items that occur together “frequently”.
- “Frequent” defined as more than “minimum support threshold”.
- Minimum support is a tunable parameter.
- Output: set of itemsets that co-occur, along with their frequencies.

Generating association rules

- Input: set of frequent itemsets; $\{\{a,b,c\}, \{a,g,f\}, \dots\}$
- Output: set of rules of the form $\{a \rightarrow bc, b \rightarrow ac, \dots\}$

Generating association rules

Given {popcorn, coke, choctop} is frequent, generate and test:

popcorn \rightarrow coke, choctop

popcorn \rightarrow coke

popcorn \rightarrow choctop

coke \rightarrow popcorn, choctop

coke \rightarrow choctop

coke \rightarrow popcorn

choctop \rightarrow coke, popcorn

choctop \rightarrow coke

choctop \rightarrow popcorn

popcorn, choctop \rightarrow coke

coke, choctop \rightarrow popcorn

coke, popcorn \rightarrow choctop

Both phases suffer
combinatorial
explosion...

Find frequent items sets with FP Growth

- Represent transaction data set as “Frequent Pattern Tree”;
- A trie, branches represent items occurring together
- Nodes contain a count
- Recursively build subtree for items to find co-occurrence & frequency.

FPGrowth: Pseudo code, initial tree build

```
For each transaction
```

```
    Count item frequencies
```

```
For each transaction
```

```
    Sort transaction by decreasing frequency, discard infrequent items
```

```
    Insert into initial FP tree
```

```
minimum_count = num_transactions * minimum_support
```

```
FPGrowth(initial FP tree, item_set=[], minimum_count)
```

FPGrowth: Pseudo code

```
FPGrowth(tree, item_set, minimum_count)
```

```
    for_each item in tree with count >= minimum_count
```

```
        conditional_tree = new FPTree()
```

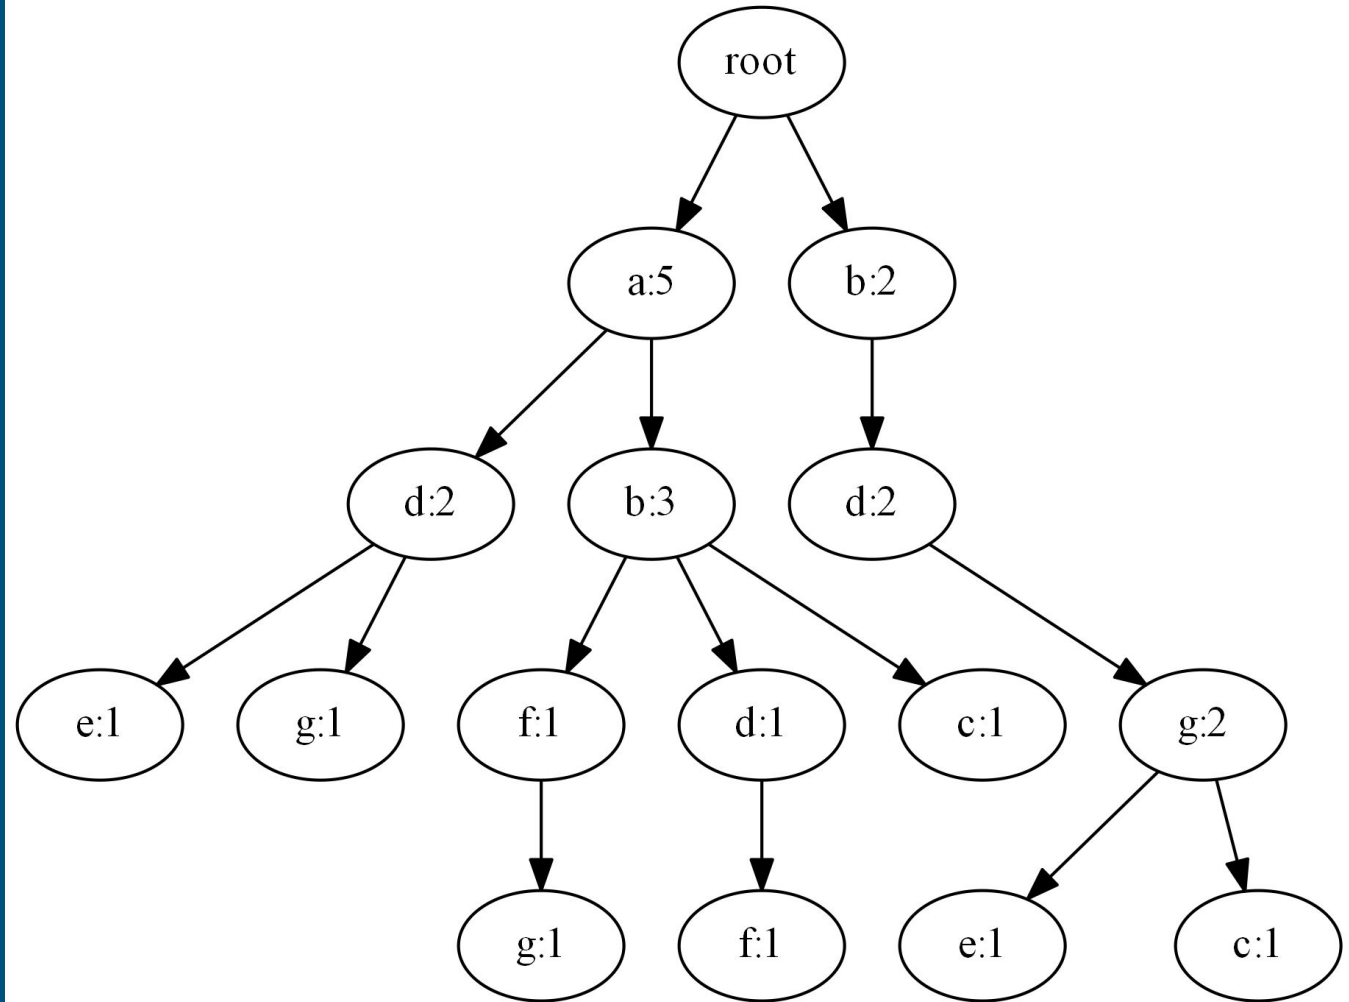
```
        for_each path in tree from item to root
```

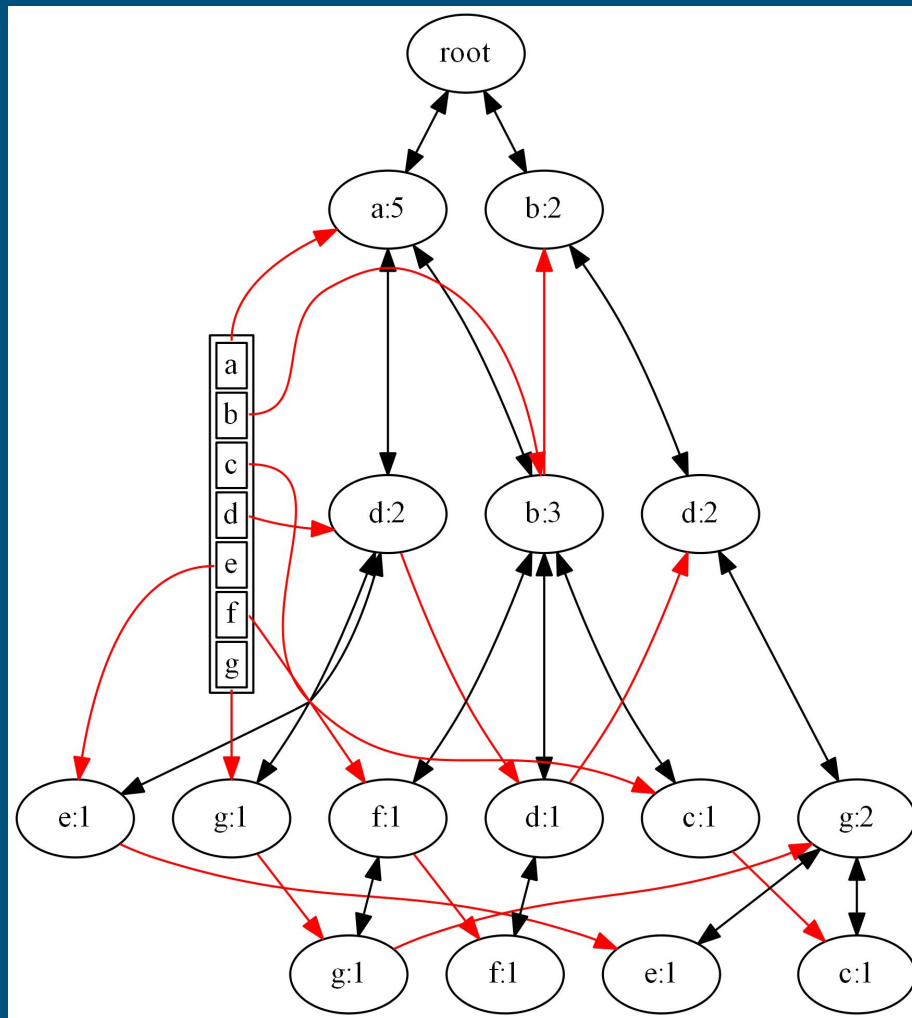
```
            Insert path excluding item into conditional_tree
```


```
        Output(item_set + item, count=conditional_tree.root.count)
```

```
        FPGrowth(conditional_tree, item_set + item, minimum_count)
```

TID	Items
1	a, d, e
2	a, b, f, g
3	a, b, d, f
4	a, b, c
5	a, d, g
6	b, c, d, g
7	b, d, e, g







Naive tree node Rust code...

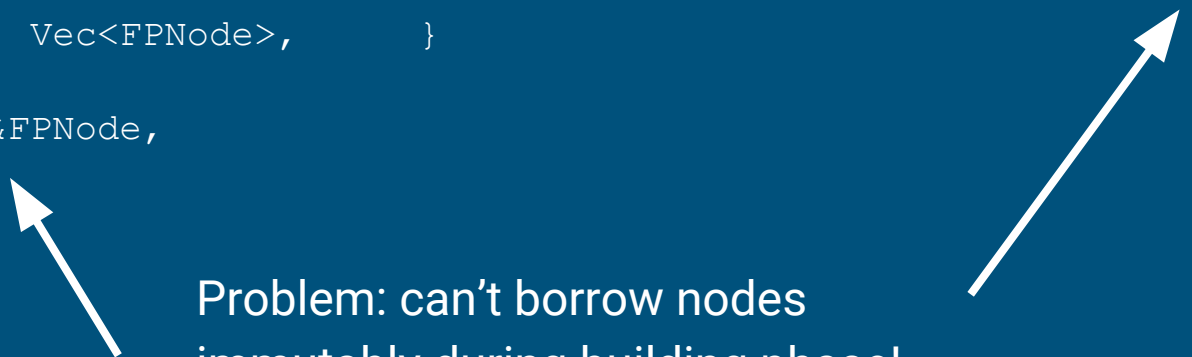
```
struct FPNode {  
    item: Item,  
  
    count: u32,  
  
    children: Vec<FPNode>,    }  
  
    parent: &FPNode,  
  
}
```

```
struct FPTree {  
    root: FPNode,  
  
    item_list: HashMap<Item, Vec<&FPNode>>,  
  
}
```

Naive tree node Rust code...

```
struct FPNode {  
    item: Item,  
    count: u32,  
    children: Vec<FPNode>,  
    parent: &FPNode,  
}  
  
struct FPTree {  
    root: FPNode,  
    item_list: HashMap<Item, Vec<&FPNode>>,  
}
```

Problem: can't borrow nodes
immutably during building phase!



Solution #1...

- Don't store reference to node's parent in node.
- Don't maintain item list.
- After building tree, traverse tree, create index nodes' parents and item list.
- Actually not terrible for performance...

Solution #2

- Store tree as a `Vec<FPNode>`.
- `FPNodes` store their parent's index.
- `FPTree` stores list of `FPNodes`' indices for each `Item`.
- Performance about the same as solution #1.
- No borrow checker shenanigans.
- Code is not as simple.

FPTree backed by an array

```
struct FPNode {  
    item: Item,  
    count: u32,  
    children: Vec<usize>,    }  
    parent: usize,  
}  
  
struct FPTree {  
    nodes: Vec<FPNode>,  
    item_list: HashMap<Item, Vec<usize>>,  
}
```

Other challenges...

Hash tables!

- Allocates lots of spare capacity
 - Reduced memory by 10X when switched nodes' children `HashMap<Item,FPNode>` to `Vec<FPNode>`!
- Very slow hashing function for primitive types.
 - FnvHash helps...
 - Replaced some `HashMap<Item, T>` with `Vec<T>`, indexed by `Item` as `usize`.
 - (May be bad idea on very big data sets)

Hash tables! (cont.)

- Use sorted `Vec<Item>` instead of `HashSet<Item>`.
 - $O(N)$ union/intersection.
 - 10% speed up by pre-calculating output `Vec<Item>` capacity!
- Can reorder `Item` indices by `ItemName` lexicographically
 - Sorted output for free!

Strings! Oh my!

- Reducing string allocations was key to improving output performance.
- Rust made it easy to create (and remove!) unnecessary allocations.
- Java implementation was actually quite fast.
 - I think due to strings & objects being passed by reference by default.

Parallelism!

Fearless concurrency

Safe Rust guarantees an absence of data races, which are defined as:

- two or more threads concurrently accessing a location of memory
- one of them is a write
- one of them is unsynchronized

<https://doc.rust-lang.org/beta/nomicon/races.html>

Concurrency


- Rust supports Go-style channels
 - “Do not communicate by sharing memory; instead, share memory by communicating.”
- Rust’s `std::threads` are OS threads
- Rust’s `std::Mutex` locks data not code

Rayon: Rust's super power

Safe foundations

- Parallel iterators make it easy to parallelize sequential code.
- Possible because everything there is Sync;
 - Can be shared across threads
- Rayon implements functional for `par_iter()`;
 - `map`, `for_each`, `filter`, `fold`

Pick the right `iter()` to
parallelize





CPU

80% 3.78 GHz



Memory

37.1/63.7 GB (58%)



Disk 0 (D:)

0%



Disk 1 (C: E:)

0%



Ethernet

S: 0 R: 0 Kbps



Ethernet

S: 24.0 R: 0 Kbps




GPU 0

NVIDIA GeForce GTX 10
1%

CPU

Intel(R) Core(TM) i9-7940X CPU @ 3.10GHz

% Utilisation



60 seconds

0

Utilisation

Speed

Base speed:

3.10 GHz

80%

3.78 GHz

Sockets:

1

Cores:

14

Processes

Threads

Handles

Logical processors:

28

253

4839

158659

Virtualisation:

Enabled

Up time

2:05:54:34

L1 cache:

896 KB

L2 cache:

14.0 MB

L3 cache:

19.3 MB



Fewer details



Open Resource Monitor

Summary

- Rust makes concurrency/parallelism easy and safe.
- Rayon is awesome.
- Rust can make the nodes in your cluster faster.
- Use Rust!

Questions?

<https://github.com/cpearce/arm-rs/>
chris@pearce.org.nz